

Major Features: Postgres 16

BRUCE MOMJIAN



POSTGRESQL is an open-source, full-featured relational database. This presentation gives an overview of the Postgres 16 release.

<https://momjian.us/presentations>



Creative Commons Attribution License


Last updated: August 2023

Postgres 16 Feature Outline

1. Logical replication from standby servers
2. Allow logical replication cycles
3. Role membership control
4. Allow indexes to use *date_trunc()*
5. Record statistics on the last sequential and index scans
6. Allow monitoring of I/O statistics
7. CPU vectorization
8. Allow libpq load balancing and control of authentication


Full item list at https://momjian.us/pgsql_docs/release-16.html and <https://www.postgresql.org/docs/16/release-16.html>.

1. Logical Replication from Standby Servers

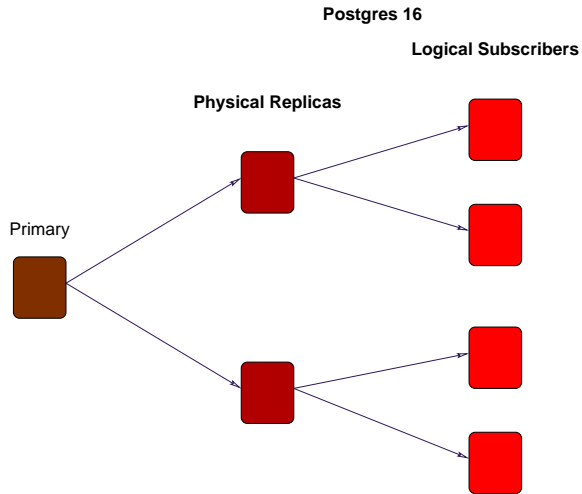


Pre-Postgres 16 Supports Only Primary-Sourced Logical Subscribers

Pre-PG 16 Logical Subscribers



Postgres 16 Supports Standby-Sourced Logical Subscribers



2. Allow Logical Replication Cycles

PG 16 Logical Subscribers on the Same Table



This is done by creating subscriptions with *slotname* = *none*, which causes only logical replication records with no slotname to be sent. This prevents records that arrived from other servers from being sent. There is no conflict resolution.

3. Role Membership Control

Role membership give two advantages:

- The ability to automatically INHERIT permissions of member roles
- The ability to execute queries as the member role via SET ROLE
- <https://www.postgresql.org/docs/current/role-membership.html>

Postgres 16 adds:

- Allow roles to INHERIT from only some members
- Allow members to be added to roles without giving them SET ROLE permission via SET ROLE FALSE
- Allow members to be added to roles if they have ADMIN permission on the role; previously create role permission was required
- Allow users who create roles to be automatically given INHERIT or SET ROLE abilities to the new role via *createrole_self_grant*
- Add *psql \drg* command to show role members and their INHERIT and SET ROLE permissions

Inheritance in Pre-Postgres 16

```
-- create role
CREATE ROLE a1;
CREATE ROLE b INHERIT;
GRANT a1 TO b;

-- create table
GRANT CREATE ON SCHEMA public TO a1;
SET ROLE a1;
CREATE TABLE a1_test (x INTEGER);
RESET ROLE;

-- SELECT table as 'b'
SET ROLE b;
SELECT * FROM a1_test;
  x
---

RESET ROLE;

-- change role inheritance status
ALTER USER b NOINHERIT;

-- SELECT table as 'b'
SET ROLE b;
SELECT * FROM a1_test;
ERROR: permission denied for table a1_test
RESET ROLE;
```


Inheritance in Posgres 16

```
-- create role
CREATE ROLE a1;
CREATE ROLE b INHERIT;
GRANT a1 TO b;

-- create table
GRANT CREATE ON SCHEMA public TO a1;
SET ROLE a1;
CREATE TABLE a1_test (x INTEGER);
RESET ROLE;

-- SELECT table as 'b'
SET ROLE b;
SELECT * FROM a1_test;
x
---
```

RESET ROLE;

```
-- change role inheritance status
ALTER USER b NOINHERIT;

-- SELECT table as 'b'
SET ROLE b;
SELECT * FROM a1_test;
x
---
```

RESET ROLE;

Inheritance in Postgres 16

```
-- create roles
```

```
CREATE ROLE a2;
```

```
CREATE ROLE a3;
```

```
CREATE ROLE a4;
```

```
-- grant membership
```

```
GRANT a2 TO b;
```

```
GRANT a3 TO b WITH INHERIT TRUE;
```

```
GRANT a4 TO b WITH INHERIT FALSE;
```

```
\drg
```

List of role grants

Role name	Member of	Options	Grantor
b	a1	INHERIT, SET	postgres
b	a2	SET	postgres
b	a3	INHERIT, SET	postgres
b	a4	SET	postgres

4. Allow Indexes to Use *date_trunc()*

```
CREATE TABLE trunc_test (x TIMESTAMP WITH TIME ZONE);
```

```
CREATE INDEX i_trunc_test ON trunc_test (date_trunc('month', x));
```

ERROR: functions in index expression must be marked IMMUTABLE

```
CREATE INDEX i_trunc_test ON trunc_test (date_trunc('month', x, 'America/New_York'));
```

```
\d trunc_test
```

Table "public.trunc_test"				
Column	Type	Collation	Nullable	Default
x	timestamp with time zone			

Indexes:

```
"i_trunc_test" btree (date_trunc('month'::text, x, 'America/New_York'::text))
```

5. Record Statistics on the Last Sequential and Index Scans

```
\d pg_stat_user_tables
```

View "pg_catalog.pg_stat_user_tables"				
Column	Type	Collation	Nullable	Default
relid	oid			
schemaname	name			
relname	name			
seq_scan	bigint			
last_seq_scan	timestamp with time zone			
seq_tup_read	bigint			
idx_scan	bigint			
last_idx_scan	timestamp with time zone			
idx_tup_fetch	bigint			
n_tup_ins	bigint			
n_tup_upd	bigint			
n_tup_del	bigint			
n_tup_hot_upd	bigint			
n_tup_newpage_upd	bigint			
n_live_tup	bigint			
n_dead_tup	bigint			
n_mod_since_analyze	bigint			
n_ins_since_vacuum	bigint			
last_vacuum	timestamp with time zone			
last_autovacuum	timestamp with time zone			
last_analyze	timestamp with time zone			
last_autoanalyze	timestamp with time zone			
vacuum_count	bigint			
autovacuum_count	bigint			
analyze_count	bigint			
autoanalyze_count	bigint			

6. Allow Monitoring of I/O Statistics

\d pg_stat_io

View "pg_catalog.pg_stat_io"				
Column	Type	Collation	Nullable	Default
backend_type	text			
object	text			
context	text			
reads	bigint			
read_time	double precision			
writes	bigint			
write_time	double precision			
writebacks	bigint			
writeback_time	double precision			
extends	bigint			
extend_time	double precision			
op_bytes	bigint			
hits	bigint			
evictions	bigint			
reuses	bigint			
fsyncs	bigint			
fsync_time	double precision			
stats_reset	timestamp with time zone			

Types of Recorded Statistics

```
SELECT backend_type, object, context FROM pg_stat_io;
```

backend_type	object	context
autovacuum launcher	relation	bulkread
autovacuum launcher	relation	normal
autovacuum worker	relation	bulkread
autovacuum worker	relation	normal
autovacuum worker	relation	vacuum
client backend	relation	bulkread
client backend	relation	bulkwrite
client backend	relation	normal
client backend	relation	vacuum
client backend	temp relation	normal
background worker	relation	bulkread
background worker	relation	bulkwrite
background worker	relation	normal
background worker	relation	vacuum
background worker	temp relation	normal
background writer	relation	normal
checkpointer	relation	normal
standalone backend	relation	bulkread
standalone backend	relation	bulkwrite
standalone backend	relation	normal
standalone backend	relation	vacuum
startup	relation	bulkread
startup	relation	bulkwrite
startup	relation	normal
startup	relation	vacuum
walsender	relation	bulkread
walsender	relation	bulkwrite

...

Sample Statistics

```
SELECT *
FROM pg_stat_io
WHERE backend_type = 'client backend' AND
      object = 'relation' AND
      context = 'normal';
-[ RECORD 1 ]-----+-----
backend_type | client backend
object       | relation
context      | normal
reads        | 3705962
read_time    | 0
writes       | 2216017
write_time   | 0
writebacks   | 0
writeback_time | 0
extends      | 20125
extend_time  | 0
op_bytes     | 8192
hits         | 81950274
evictions    | 3714568
reuses       | (null)
fsyncs       | 0
fsync_time   | 0
stats_reset  | 2023-08-23 21:30:36.252786-04
```


7. CPU Vectorization

- Allow JSON string parsing to use vector operations
- Add support for SSE2 (Streaming SIMD Extensions 2) vector operations on x86-64 architectures
- Add support for Advanced SIMD (Single Instruction Multiple Data) (NEON) instructions on ARM architectures

8. Allow libpq Load Balancing and Control of Authentication

- libpq connection option *load_balance_hosts=random* now randomly chooses a listed host
 - combined with *target_session_attrs=standby*, this allows load balancing among standby servers
- libpq option *require_auth* allows client to specify acceptable authentication methods
 - for example, *require_auth=scram-sha-256* prevents other password methods from being used

Conclusion



<https://momjian.us/presentations>

<https://www.flickr.com/photos/thomasletholsen/>